

Visualization (Data Transformation)

Peter Ganong and Maggie Shi

January 21, 2026

Introduction: roadmap

- Putting this lecture in context
- Introducing the `movies` dataset
 - `load data`
 - `shape`
 - `head()`

Putting this lecture in context

- This lecture explores methods for *transforming* data, focusing on aggregation.
 - We will be mostly following Chapter 3 in the data visualization (Heer et al.) book
- Fundamental problem in data visualization: in most cases, **you do not want to show every single data point** in your dataset.
- Instead, you want to extract patterns which you (the analyst) think are interesting.

Aggregation

- One nice thing about `altair` is that it nudges you to aggregate.
- One example: if you try to make a plot with 10,000 dots, it will give you an error: `MaxRowsError: The number of rows in your dataset is greater than the maximum allowed (500).`
 - Help file: “This is not because Altair cannot handle larger datasets, but it is because it is important for the user to think carefully about how large datasets are handled.”
 - More details [here](#)

Load packages and data

```
1 import pandas as pd
2 import altair as alt
3
4 movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'
5 movies = pd.read_json(movies_url)
```

```
[{"Title": "The Land Girls", "US_Gross": 146083, "Worldwide_Gross": 146083, "US_DVD_Sales": null, "Production_Budget": 8000000, "Release_Date": "Jun 12 1998", "MPAA_Rating": "R", "Running_Time_min": null, "Distributor": "Gramercy", "Source": null, "Major_Genre": null, "Creative_Type": null, "Director": null, "Rotten_Tomatoes_Rating": null, "IMDB_Rating": 6.1, "IMDB_Votes": 1071}, {"Title": "First Love, Last Rites", "US_Gross": 10876, "Worldwide_Gross": 10876, "US_DVD_Sales": null, "Production_Budget": 300000, "Release_Date": "Aug 07 1998", "MPAA_Rating": "R", "Running_Time_min": null, "Distributor": "Strand", "Source": null, "Major_Genre": "Drama", "Creative_Type": null, "Director": null, "Rotten_Tomatoes_Rating": null, "IMDB_Rating": 6.9, "IMDB_Votes": 207}]
```

An aside on JSON

- Movies database is stored as a `.json` at [this URL](#)
- Recall `altair` that writes Vega-lite, which is also recorded in JSON!
 - JSON is just a “syntax” to store text, numbers, etc. in a human-readable way
 - In spatial lectures, we will also encounter the “geojson” format, which can store geographic features

head()

```
1 movies.head(5)
```

	Title	US_Gross	Worldwide_Gross	US_DVD_Sales	Production_Budget	Release_Date	MPAA_Rating	Runni
0	The Land Girls	146083.0	146083.0	NaN	8000000.0	Jun 12 1998	R	NaN
1	First Love, Last Rites	10876.0	10876.0	NaN	300000.0	Aug 07 1998	R	NaN
2	I Married a Strange Person	203134.0	203134.0	NaN	250000.0	Aug 28 1998	None	NaN
3	Let's Talk About Sex	373615.0	373615.0	NaN	300000.0	Sep 11 1998	None	NaN
4	Slam	1009819.0	1087521.0	NaN	1000000.0	Oct 09 1998	R	NaN

shape

```
1 movies.shape
```

```
(3201, 16)
```


With 3201 movies, we are going to need to do some transformation if we want to uncover any patterns in the data!

Variables of interest

- **Rotten Tomatoes** ratings are determined by taking “thumbs up” and “thumbs down” judgments from film critics and calculating the percentage of positive reviews.
- **IMDB ratings** are formed by averaging scores (ranging from 1 to 10) provided by the site’s users.

Exploring the raw data

```
1 alt.Chart(movies_url).mark_circle().encode(  
2     alt.X('Rotten_Tomatoes_Rating:Q', title = "Rotten Tomatoes Rating (%)")  
3     alt.Y('IMDB_Rating:Q', title = "IMDB Rating")  
4 )
```



Recall from last lecture: label when scale is %!

Aggregation

Aggregation: roadmap


In previous lectures, we actually already saw aggregation via `average()` and `min()`. We just didn't talk explicitly about that step. Now, we examine it more carefully.

- `average()`
- interquartile range
- do-pair-share

The Altair documentation includes the [full set of available aggregation functions](#).

average()

```
1 alt.Chart(movies_url).mark_bar().encode(  
2     alt.X('average(Rotten_Tomatoes_Rating):Q', title = "Average Rotten Tomatoes Rating (%)",  
3     alt.Y('Major_Genre:N', title = "Genre")  
4 )
```




- This plot is fine, but hard to interpret takeaways quickly.
- Discussion Question: what does the y-axis seem to be sorted on? Why?

average() with sort(...)

More useful: sort the bars vertically, based on x-axis encoding

```
1 alt.Chart(movies_url).mark_bar().encode(
2     alt.X('average(Rotten_Tomatoes_Rating):Q', title = "Average Rotten Tomatoes Rating (%)"),
3     alt.Y('Major_Genre:N', title = "Genre",
4         sort=alt.EncodingSortField(op='average', field='Rotten_Tomatoes_Rating', order='descending')
5     )
6 )
```




This focuses the viewer's attention on which movie types are most and least popular

Interquartile range

Plot 1st and 3rd quartiles, then sort by median.

```
1 alt.Chart(movies_url).mark_bar().encode(
2     alt.X('q1(Rotten_Tomatoes_Rating):Q', title = "Rotten Tomatoes Rating (%)"),
3     alt.X2('q3(Rotten_Tomatoes_Rating):Q'),
4     alt.Y('Major_Genre:N', sort=alt.EncodingSortField(op='median', field='Rotten_Tomatoes_Rati
5         title = "Genre"
6     )
7 )
```



Discussion question: what can you learn from the IQR plot that you could not learn from the plot with just `average()`?

Aggregation functions

- Distribution: `min()`, `q1()`, `median()`, `mean()`, `q3()`, `max()`
- Dispersion: `variance()`, `stdev()`
- Bootstrap confidence intervals: `ci0()`, `ci1()`
- Full list [here](#)


Case study: when are the highest grossing films?

```
1 movies_gross = movies[['US_Gross', 'Release_Date']]  
2 movies_gross.head()
```

	US_Gross	Release_Date
0	146083.0	Jun 12 1998
1	10876.0	Aug 07 1998
2	203134.0	Aug 28 1998
3	373615.0	Sep 11 1998
4	1009819.0	Oct 09 1998

A first pass

```
1 alt.Chart(movies_url).mark_point().encode(  
2     alt.X('Release_Date:T', title = "Release Date"),  
3     alt.Y('US_Gross:Q', title = "US Gross ($)")  
4 )
```



Obviously we need to aggregate.

Also: what bug in the data does this plot reveal?

Do-pair-share

1. *Do* – make a plot on your own
 2. *Pair* – compare your results with person next to you
 3. *Share* – discuss results as a class
- **Question:** What time of year are the highest grossing films released? Aggregate both the x- and the y-variables.
 - There are several ways to approach answering this question. What seems most reasonable to you?

Data aggregation: Do-pair-share

Starter code in lecture [dps_highest_grossing_film.qmd](#) file:

```
1 import pandas as pd
2 import altair as alt
3 movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'
4 movies = pd.read_json(movies_url)
5
6 # unaggregated scatter plot
7 alt.Chart(movies_url).mark_point().encode(
8     alt.X('Release_Date:T', title = "Release Date"),
9     alt.Y('US_Gross:Q', title = "US Gross ($"))
10 )
```

Documentation on working with time units [here](#)

More on time units

Temporal variables can be transformed into a variety of other time units

- `year`
- `quarter`
- `month`
- `date` (numeric day in month)
- `day` (day of the week)
- `hours`
- `yearmonth`
- `hoursminutes`

Aggregation: summary

- Many built-in aggregation functions to quantify distribution, dispersion, and characterize data
- Dates: see prior slide

Advanced data transformation

Advanced data transformation: introduction

- Two ways to aggregate data in `altair`
 - Within the encoding itself:
`alt.Y('median(US_Gross):Q')`
 - Separately using a top-level aggregate transform
- Doing it in the encoding is fine for simple transformations
- But for advanced transformations, we'll have to define it separately

Advanced data transformation: roadmap

- `transform_calculate()`
- `transform_timeunit()`
- `transform_filter()`
- do-pair-share
- `transform_aggregate()`
- `transform_window()`

These are all written in the [Vega expression language](#).

Connection to pandas operations

One way to think of these verbs is that they are fundamental to any data analysis project and so in any/every package you learn, you need to know how to do these.

Purpose	Vega	pandas equivalent
Define a new variable	<code>transform_calculate()</code>	<code>df['new_col']</code>
Filter to subset of rows	<code>transform_filter(cond)</code>	<code>df.loc[cond]</code>
Aggregate function - collapse number of rows down to one per group	<code>transform_aggregate(groupby(...))</code>	<code>df.groupby('A').agg('mean')</code>
Window function - transform across multiple rows, keeps same num. of rows)	<code>transform_window(sum())</code>	<code>df['values'].cumsum()</code>

Connection to `pandas` operations

- You already know how to do these all in `pandas` so it is not conceptually new.
- Why bother doing it in `altair`?
 - Exploratory data analysis can be done faster in `altair`: manipulate data and plot simultaneously
 - Aggregation and transformations are **temporary** – don't need to define and keep track of new aggregated dataframes

transform_calculate and transform_timeunit

- `transform_calculate()` uses expressions for writing basic formulas
 - Math functions: `min()`, `random()`, `round()`
 - Statistical functions: `sampleNormal()`, `sampleUniform()`
 - String functions: `length()`, `lower()`, `substring()`
- Use `transform_timeunit()` when working with Temporal variables
 - `month()`, `quarter()`, `yearmonth()`

transform_calculate case study

Question: what time of year do US movies make money abroad?

```
1 alt.Chart(movies_url).mark_area().transform_calculate(  
2     NonUS_Gross='datum.WORLDWIDE_GROSS - datum.US_GROSS'  
3 ).encode(  
4     alt.X('month(Release_Date):T', title = "Release Month"),  
5     alt.Y('median(NonUS_Gross):Q', title = "Median Non-US Gross ($)")  
6 )
```

transform_calculate case study

Question: what time of year do US movies make money abroad?

```
1 alt.Chart(movies_url).mark_area().transform_calculate(  
2     NonUS_Gross='datum.WORLDWIDE_GROSS - datum.US_GROSS'  
3 ).encode(  
4     alt.X('month(Release_Date):T', title = "Release Month"),  
5     alt.Y('median(NonUS_Gross):Q', title = "Median Non-US Gross ($)")  
6 )
```

- `NonUS_Gross` is a variable we're defining *temporarily*

transform_calculate case study

Question: what time of year do US movies make money abroad?

```
1 alt.Chart(movies_url).mark_area().transform_calculate(  
2     NonUS_Gross='datum.WORLDWIDE_GROSS - datum.US_GROSS'  
3 ).encode(  
4     alt.X('month(Release_Date):T', title = "Release Month"),  
5     alt.Y('median(NonUS_Gross):Q', title = "Median Non-US Gross ($)")  
6 )
```

- `datum` is how you reference the underlying dataset within a transformation expression
- Here, `datum` means `movies_url`

transform_calculate case study


Question: what time of year do US movies make money abroad?

```
1 alt.Chart(movies_url).mark_area().transform_calculate(  
2     NonUS_Gross='datum.WORLDWIDE_GROSS - datum.US_GROSS'  
3 ).encode(  
4     alt.X('month(Release_Date):T', title = "Release Month"),  
5     alt.Y('median(NonUS_Gross):Q', title = "Median Non-US Gross ($)")  
6 )
```

- After defining `NonUS_Gross`, we can use like any other variable within `movies_url`
- It can be combined with other aggregation methods

transform_calculate case study

Question: what time of year do US movies make money abroad?



transform_filter

- Goal: show just movies before 1970


```
1 alt.Chart(movies_url).mark_circle().transform_filter(  
2     'year(datum.Release_Date) < 1970').encode(  
3     alt.X('Rotten_Tomatoes_Rating:Q', title = "Rotten Tomatoes Rating (%)")  
4     alt.Y('IMDB_Rating:Q', title = "IMDB Rating"),  
5 ).properties(title = "Pre-1970 Sample")
```

- `transform_filter` filters the dataset based on an expression
- Like `transform_aggregate`, this filtering is *temporary*

transform_filter

- Goal: show just movies before 1970

```
1 alt.Chart(movies_url).mark_circle().transform_filter(  
2     'year(datum.Release_Date) < 1970').encode(  
3         alt.X('Rotten_Tomatoes_Rating:Q', title = "Rotten Tomatoes Rating (%)")  
4         alt.Y('IMDB_Rating:Q', title = "IMDB Rating")  
5     ).properties(title = "Pre-1970 Sample")
```



Do-pair-share

- Make two plots that compare ratings before and after 1970
 - Use `transform_filter()` to create a plot of ratings before vs. after 1970, then append side-by-side
 - Plot before and after 1970 on one plot
 - Use `transform_aggregate()` to create a categorical variable to indicate whether an observation is from before or after 1970.
 - Encode the color of the mark depending on the value of that categorical variable
- These plots show equivalent information. Which do you prefer and why?

Do-pair-share

Starter code in lecture `dps_1970.qmd` file:

```
1 import pandas as pd
2 import altair as alt
3 movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'
4 movies = pd.read_json(movies_url)
5
6 # scatter plot, filtered to < 1970
7 alt.Chart(movies_url).mark_circle().encode(
8     alt.X('Rotten_Tomatoes_Rating:Q', title = "Rotten Tomatoes Rating (%)" ),
9     alt.Y('IMDB_Rating:Q', title = "IMDB Rating")
10 ).transform_filter('year(datum.Release_Date) < 1970'
11 ).properties(title = "Pre-1970 Sample")
```

Hint: recall `graphA | graphB` plots `graphA` next to `graphB`

transform_window: case study


Question: who are the top grossing directors of all time?

```
1 alt.Chart(movies_url).mark_bar().transform_aggregate(  
2     Gross='sum(Worldwide_Gross)',  
3     groupby=['Director']  
4 ).encode(  
5     alt.X('Gross:Q', title = "Worldwide Gross ($)'),  
6     alt.Y('Director:N', sort=alt.EncodingSortField(  
7         op='max', field='Gross', order='descending'  
8     ),  
9     title = "Director")  
10 )
```

- First, sum `Worldwide_Gross` for each director to make `Gross`
- Then, plot in descending order of `Gross`

transform_window: case study

Question: who are the top grossing directors of all time?



transform_window: case study

That's a lot of directors! Let's restrict to the **top 10**

```
1 alt.Chart(movies_url).mark_bar().transform_aggregate(  
2     Gross='sum(Worldwide_Gross)',  
3     groupby=['Director']  
4 ).transform_window(  
5     Rank='rank()',  
6     sort=[alt.SortField('Gross', order='descending')]  
7 ).transform_filter(  
8     'datum.Rank <= 10'  
9 ).encode(  
10     alt.X('Gross:Q', title = "Worldwide Gross ($)'),  
11     alt.Y('Director:N', sort=alt.EncodingSortField(  
12         op='max', field='Gross', order='descending'  
13     ), title = "Director")  
14 )
```

- Use `transform_window()` when we want to add a variable and keep the *same* number of rows
- Opposed to `transform_aggregate()` which *collapses* the number of rows


transform_window: case study

That's a lot of directors! Let's restrict to the **top 10**

```
1 alt.Chart(movies_url).mark_bar().transform_aggregate(  
2     Gross='sum(Worldwide_Gross)',  
3     groupby=['Director']  
4 ).transform_window(  
5     Rank='rank()',  
6     sort=[alt.SortField('Gross', order='descending')]  
7 ).transform_filter(  
8     'datum.Rank <= 10'  
9 ).encode(  
10     alt.X('Gross:Q', title = "Worldwide Gross ($)'),  
11     alt.Y('Director:N', sort=alt.EncodingSortField(  
12         op='max', field='Gross', order='descending'  
13     ), title = "Director")  
14 )
```

After ranking, use `transform_filter()` to restrict to the top 10 highest-ranked

transform_window: case study




- `null` is not a director, and we certainly don't want to say they're the highest-grossing director.
- So let's remove that -> `transform_filter()` again

transform_window: case study

```
1 alt.Chart(movies_url).mark_bar().transform_aggregate(  
2     Gross='sum(Worldwide_Gross)',  
3     groupby=['Director']  
4 ).transform_window(  
5     Rank='rank()',  
6     sort=[alt.SortField('Gross', order='descending')]  
7 ).transform_filter(  
8     'datum.Rank <= 11'  
9 ).transform_filter(  
10    'datum.Director != null'  
11 ).encode(  
12     alt.X('Gross:Q', title = "Worldwide Gross ($)'),  
13     alt.Y('Director:N', sort=alt.EncodingSortField(  
14         op='max', field='Gross', order='descending'  
15     ), title = "Director")  
16 )
```

Note that the `transform_filter()` is now `<=11`. Why?

transform_window: case study



Steven Spielberg has been quite successful in his career!

transform_window: do-pair-share

- Showing sums might favor directors who have had longer careers, and so have made more movies and thus more money.
- What happens if we change the choice of aggregate operation?
- Who is the most successful director in terms of average gross per film?

transform_window: do-pair-share

Starter code in `dps_directors.qmd`:

```
1 import pandas as pd
2 import altair as alt
3 movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'
4 movies = pd.read_json(movies_url)
5
6 alt.Chart(movies_url).mark_bar().transform_filter(
7     'datum.Director != null'
8 ).transform_aggregate(
9     Gross='sum(Worldwide_Gross)',
10    groupby=['Director']
11 ).transform_window(
12     Rank='rank()',
13     sort=[alt.SortField('Gross', order='descending')]
14 ).transform_filter(
15     'datum.Rank < 10'
16 ).encode(
17     alt.X('Gross:Q', title = "Worldwide Gross ($"),
18     alt.Y('Director:N'. sort=alt.EncodingSortField(
```

Advanced data transformation: summary

Purpose	Vega	pandas equivalent
Define a new variable	<code>transform_calculate()</code>	<code>df['new_col']</code>
Filter to subset of rows	<code>transform_filter(cond)</code>	<code>df.loc[cond]</code>
Aggregate function - reduces number of rows down to one per group	<code>transform_aggregate(groupby(...))</code>	<code>df.groupby('A').agg('mean')</code>
Window function - transform across multiple rows, keeps same num. of rows)	<code>transform_window(sum())</code>	<code>df['values'].cumsum()</code>

- Altair actually has 19 transformation methods (and counting...) and we have only covered four of them.
- Read about the rest of them [here](#).